CLOUD COMPUTING APPLICATIONS

Hbase Usage API

Roy Campbell & Reza Farivar

HBase: Overview

- HBase is a distributed column-oriented data store built on top of HDFS
- HBase is an Apache open source project whose goal is to provide storage for the Hadoop Distributed Computing
- Data is logically organized into tables, rows and columns

HBase:
Part of Hadoop’s Ecosystem

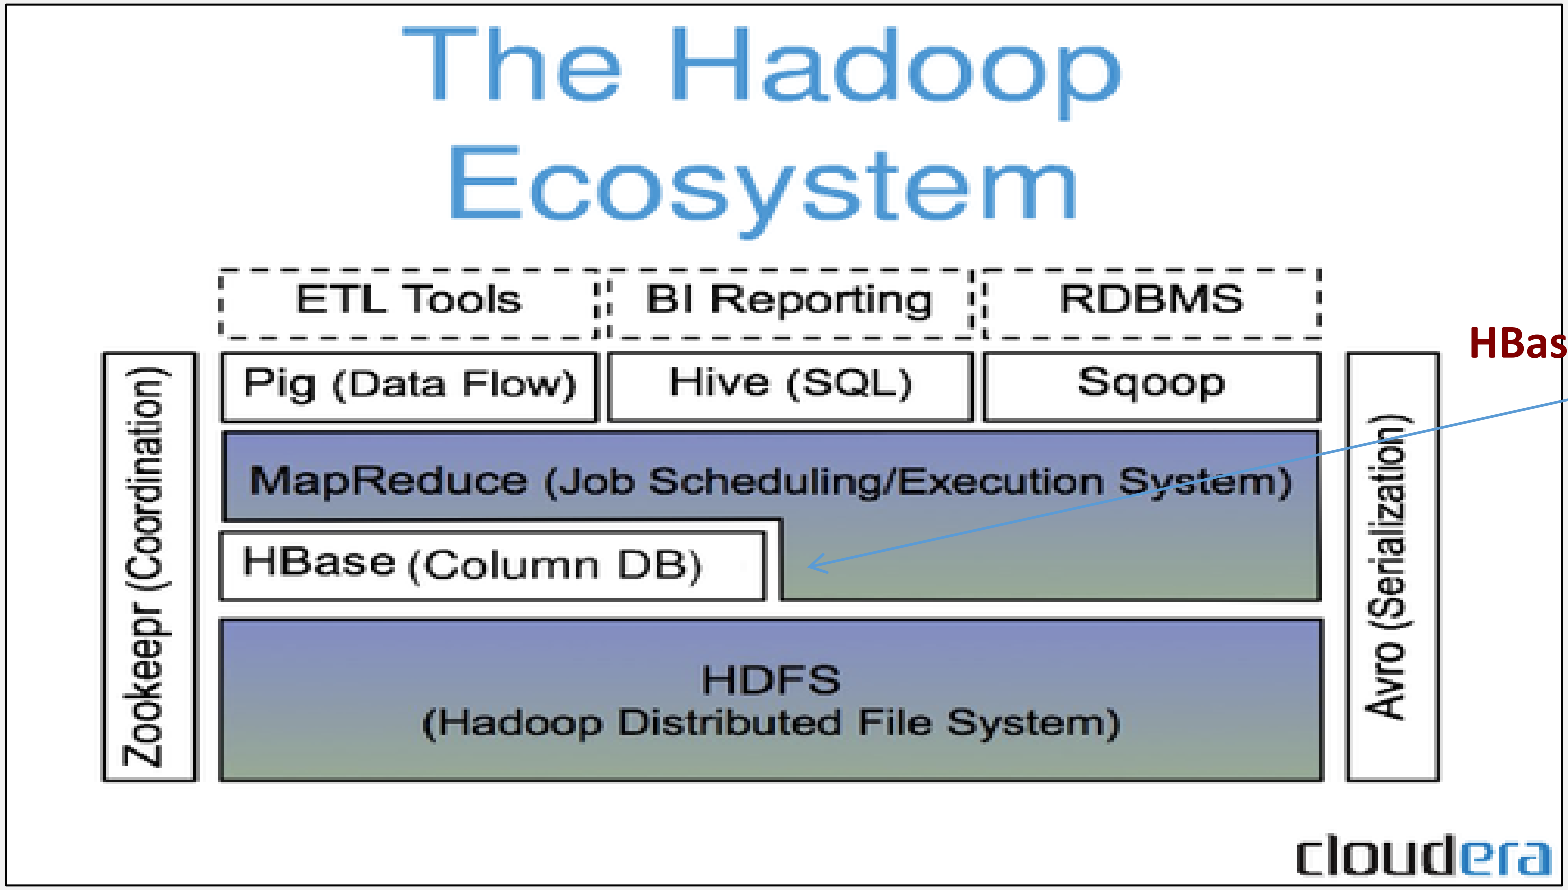

HBase is built on top of HDFS

HBase files are internally stored in HDFS

HBase vs. HDFS

- ***HDFS*** is good for batch processing (scans over big files)
  - Not good for record lookup
  - Not good for incremental addition of small batches
  - Not good for updates
- ***Hbase*** addresses the above points
  - Fast record lookup
  - Support for record-level insertion
  - Support for updates

Data Model

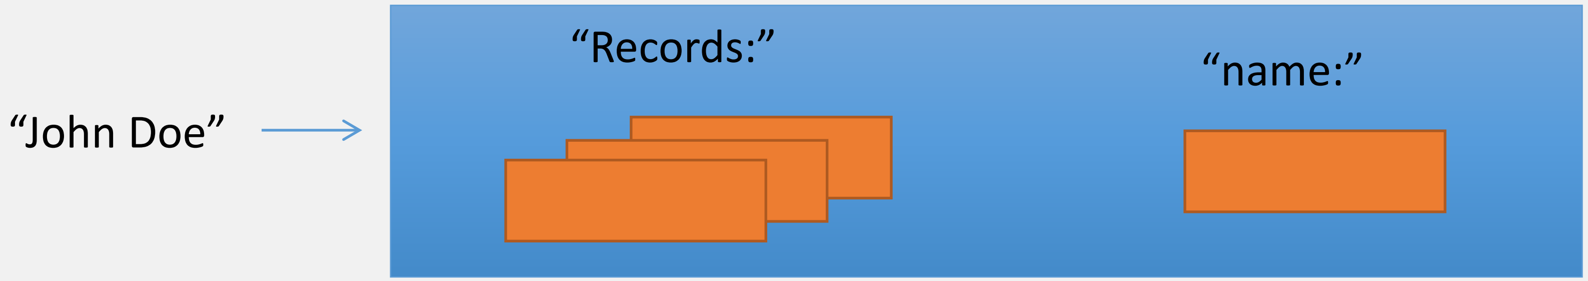

- A table in Bigtable is a sparse, distributed, persistent multidimensional sorted map
- Map indexed by a row key, column key, and a timestamp
  - (row:string, column:string, time:int64) → uninterpreted byte array
- Supports lookups, inserts, deletes
  - Single row transactions only

Notes on Data Model

- HBase schema consists of several ***Tables***
- Each table consists of a set of ***Column Families***
  - Columns are not part of the schema
- HBase has ***Dynamic Columns***
  - Because column names are encoded inside the cells
  - Different cells can have different columns

“Roles” column family has different columns in different cells

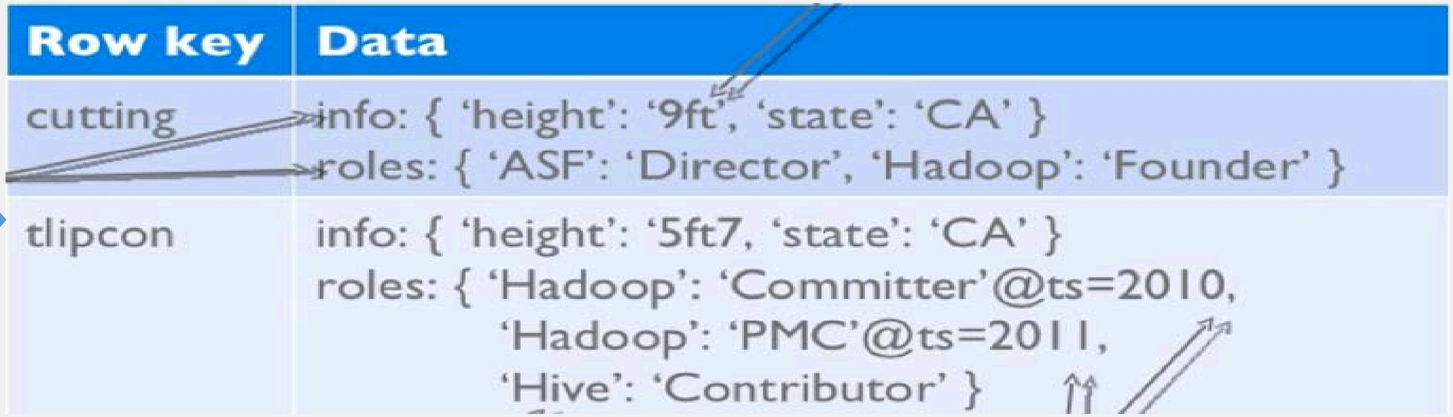

Row key	Data

cutting	info: { ‘height’: ‘9ft’, ‘state’: ‘CA’ }
	roles: { ‘ASF’: ‘Director’, ‘Hadoop’: ‘Founder’ }

tlipcon	info: { ‘height’: ‘5ft7, ‘state’: ‘CA’ }
	roles: { ‘Hadoop’: ‘Committer’@ts=2010,
	‘Hadoop’: ‘PMC’@ts=2011,
	‘Hive’: ‘Contributor’ }

Notes on Data Model

- The ***version number*** can be user-supplied
  - Even does not have to be inserted in increasing order
  - Version number are unique within each key
- Table can be very sparse
  - Many cells are empty
- ***Keys*** are indexed as the primary key

Has two columns
[cnnsi.com & my.look.ca]

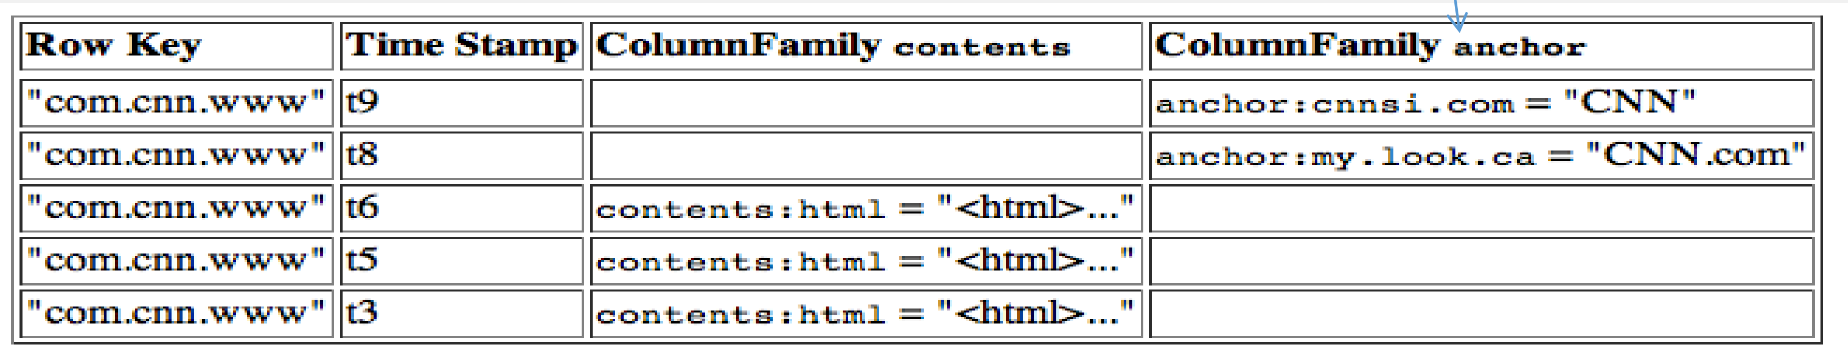

Row Key	Time Stamp	ColumnFamily contents	ColumnFamily anchor
"com.cnn.www"	t9		anchor:cnnsi.com = "CNN"
"com.cnn.www"	t8		anchor:my.look.ca = "CNN.com"
"com.cnn.www"	t6	contents:html = "<html>..."	
"com.cnn.www"	t5	contents:html = "<html>..."	
"com.cnn.www"	t3	contents:html = "<html>..."	

Rows and Columns

- Rows maintained in sorted lexicographic order
  - Applications can exploit this property for efficient row scans
  - Row ranges dynamically partitioned into tablets
- Columns grouped into column families
  - Column key = *family:qualifier*
  - Column families provide locality hints
  - Unbounded number of columns